

What Is Hermes Agent ?

Author: Nous Research @NousResearch · Feb 25 **URL:** <https://x.com/0xMovez/status/2045080054917476451>

My last article on how I built the Weather Clawdbot agent for Polymarket trading reached 2.5M views and helped many Polymarket traders and developers.

A month ago, @NousResearch released the Hermes agent, I started using it daily for research & Polymarket trading - and I can confidently say it's the best agent on the market.


It's no secret that agents and bots on Polymarket earn millions of dollars daily in areas like weather trading, crypto, and sports markets.

In these fields, the right algorithm and execution speed are everything - that's why automated, self-learning AI agents have such a strong edge. They don't sleep, they don't have emotions, and they remember everything.


Example of such agents/bots:

- ColdMath - weather trading bot that turned \$300 → \$219K in 3 months.

Profile: <https://polymarket.com/@coldmath?via=following>




- Sharky6999 - crypto trading bot with \$819K PnL and a 99,3% win-rate > Profile: <https://polymarket.com/@sharky6999?via=following>



- RN1 - best Polymarket sports bot, turned a \$1.2K deposit into \$7.3M

Profile: <https://polymarket.com/@rn1?via=following>




In this article we will discover: what Hermes Agent is, why it replaced OpenClaw, step-by-step install, where to find 700+ skills, and a complete walkthrough of building a Polymarket weather trading agent.

Lets start !

What Is Hermes Agent ?



Hermes Agent is an open-source, self-hosted AI agent built by Nous Research - the team behind YaRN, Nomos, and the Psyche model families, that was released on February 25, 2026.




Here's the simplest way to understand it.

When you hear "AI," your brain imagines a chatbot. Hermes is not a chatbot. It's more like hiring a staff member who never sleeps, remembers everything you've ever told them, gets better at their job every single day, and is always available through your phone - even when your computer is off.

3 things that make Hermes different:

Hermes through a three-layer model that resonated widely in the community - and it's the clearest way to understand what Hermes actually does for you.



screenshot is taken from @blockmates article about Hermes agent


Knowledge Layer: Built-in memory, session search, LLM-Wiki skill, optional Honcho integration. Agent doesn't just answer - it accumulates knowledge over time

Execution Layer: Multi-agent profiles, child agents, tool system, MCP support, persistent machine access. Agent doesn't just respond - it decomposes tasks, runs them in parallel, and delegates

Output Layer: Cron jobs, gateway delivery to Telegram/Slack/Discord, Web UI, file outputResults flow back into your real workflow - not trapped in a chat window.

Here are the top 3 things that make it the #1 agent in the market for me:

1. Persistent Memory



Hermes doesn't forget you when you close the tab. It stores two core memory files: MEMORY.md (environment facts, conventions, experiences) and USER.md (your preferences, communication style, expectations).

These are loaded as a frozen snapshot into every new session. It also stores full conversation history in a SQLite database with full-text search. Past conversations are searchable, summarizable, and retrievable on demand.

2. Self-Improving Skills

SKILLS

40+ bundled skills covering MLOps, GitHub workflows, research, and more. The agent creates new skills on the fly and shares them via the [open agentskills.io](https://agentskills.io) format.

This is the killer feature. After completing a complex task (roughly 5+ tool calls), Hermes automatically creates a "skill" -a structured markdown file that captures the procedure, known pitfalls, and verification steps.

Next time it encounters a similar task, it loads that skill and executes faster, better, and the way you prefer. The longer you run Hermes, the smarter it gets at your specific workflows.

3. Always-On Execution


LIVES WHERE YOU DO

Telegram, Discord, Slack, WhatsApp, Signal, Email, CLI – and a growing list of platforms. Start on one, pick up on another.

Hermes runs on your server 24/7. It connects to Telegram, Discord, Slack, WhatsApp, Signal, Email, and 15+ other platforms through a single gateway process.

You message it whenever. It has built-in cron scheduling - you can write natural-language schedules like "Every morning at 8am, scan these GitHub repos and summarize changes to my Telegram." It runs unattended.

The practical impact is of those features:



As @0xJeff described in his article after three weeks of using Hermes as a personal analyst: instead of opening Telegram, X, Rabby wallet, and Coingecko every morning.

He just open Discord, check what Hermes has prepared, give feedback, and the agent improves. Total cost: \$5-\$10/month.

Why Hermes is winning over OpenClaw

OPENCLAW	VS	HERMES
<i>The Tool Runner</i>		<i>The Meta-Agent</i>
• Config-first, channel-first architecture • TypeScript / Node runtime • 145K+ GitHub stars • 13,000+ skills on ClawhHub • 50+ messaging platforms • Session-based memory (manual setup) • No built-in learning loop • Skills installed & maintained manually • 9 CVEs disclosed in 4 days (March 2026) • Creator joined OpenAI — now community-led		• Learning-loop-first architecture • Python runtime + persistent memory • 64K+ GitHub stars in 2 months • Creates its OWN skills from experience • 15+ platforms — growing weekly • Built-in bounded, curated memory (MEMORY.md + USER.md) • Honcho dialectic user modeling across sessions • Built-in Atropos RL training pipeline • Zero agent-specific CVEs • Actively shipped by Nous Research — 1 major release/week

OpenClaw (formerly Clawdbot, then Moltbot) was the breakout AI agent of early 2026. It hit 145,000+ GitHub stars, gained 20,000 stars in a single 24-hour period, and even drove Mac Mini sales to sell-out levels. It was the project that proved personal AI agents are possible and desirable.

But users who've tried both are migrating to Hermes. Here is two mains points for me:

1. Learning Loop: fundamental difference

OpenClaw is a static agent framework.

You install skills manually. They don't improve. Your agent on day 30 performs identically to your agent on day 1 -unless you manually update its configuration.

Hermes has a closed learning loop.

Every ~15 tool calls, it pauses, reviews what worked and what failed, and writes a reusable skill. These skills aren't hidden - they're readable, editable markdown files in ~/.hermes/skills/. You can review them, tweak them, delete bad ones. But the default behavior is: the agent gets measurably better with use.

2. Memory: Bounded vs. Manual

OpenClaw's memory between sessions requires manual setup and third-party tools (QMD or similar augmentation). Without these add-ons, each session starts with limited context from

prior work.

Hermes has built-in bounded, curated memory. The key word is bounded - Hermes doesn't blindly stuff everything into context. It keeps high-value, low-change information persistent, and retrieves historical sessions on demand via full-text search. It also actively manages its own memory, including intelligent context compression during long sessions.

Complete comparison openclaw/hermes

FEATURE	OPENCLAW	HERMES AGENT
SELF-IMPROVING SKILLS	✗ Manual install only	✓ Auto-creates & self-improves
PERSISTENT MEMORY	🔌 Requires third-party (QMD)	✓ Built-in + Honcho option
LEARNING LOOP	✗ Static — day 1 = day 30	✓ Every ~15 tool calls
MULTI-AGENT	🔌 Manual multi-instance	✓ Native child agents, isolated context
MESSAGING PLATFORMS	50+ — broadest coverage	15+ — growing fast
SKILL ECOSYSTEM	13,000+ on ClawHub	700+ on agentskills.io (4 weeks old)
SECURITY RECORD	9 CVEs in 4 days + 135K exposed instances	0 CVEs + 65-rule security scanner
CRON / SCHEDULING	✓ Built-in	✓ Natural language + isolated sessions
BROWSER AUTOMATION	✓ Standard headless	✓ Camoufox anti-detection stealth
MCP PROTOCOL	✓	✓ Native since v0.8.0
MIGRATION PATH	—	✓ hermes claw migrate
RL TRAINING	✗	✓ Atropos integration built-in
RUNTIME	TypeScript / Node	Python 3.11+
LICENSE	Open Source	MIT (full commercial freedom)
COST TO RUN	\$0 software + LLM API	\$0 software + LLM API (\$5/mo VPS)

OpenClaw proved the concept. Hermes perfected the architecture. If you need 50+ platform integrations and don't care about self-improvement, OpenClaw still has broader coverage.


But if you want an agent that compounds in capability, has stronger security, and gets smarter the more you use it - Hermes is the clear choice in April 2026.

How to install hermes agent (guide)

The install takes under 5 minutes. Hermes supports Linux, macOS, and WSL2. Native Windows is not supported - use WSL2 (Ubuntu) if you're on Windows.

Step 1 - prepare VPS

I want my agent to run 24/7, so I will be deploying it on a Hetzner VPS server.



Create an account at hetzner.com → pass simple KYC verification and rent a VPS server (regular performance) in the admin panel.

Step 2 - connect to VPS

Run the command below in your Mac Terminal or use "Terminus" for Windows.

Step 3 - install Hermes

Run one command below. The installer handles platform-specific setup automatically.

```

‡ Hermes Agent Installer

An open source AI agent by Nous Research.

✓ Detected: linux (ubuntu)
→ Checking for uv package manager...
→ Installing uv (fast Python package manager)...
installing to /root/.local/bin
  
```

Step 4 - choose model

After installation is done, the guide will prompt you to choose a model. I recommend using ChatGPT (Codex) or Anthropic (Claude).



Step 5 - set up messaging gateway

Run command below to connects Hermes to your messaging platform - so you can talk to it from your phone.

```
◆ Messaging Platforms

— [📄] Telegram Setup —

1. Open Telegram and message @BotFather
2. Send /newbot and follow the prompts to create your bot
3. Copy the bot token BotFather gives you
4. To find your user ID: message @userinfobot - it replies with your numeric ID

Paste the token from @BotFather (step 3 above).
[ Bot token:
✓ Saved TELEGRAM_BOT_TOKEN

Paste your user ID from step 4 above.
The gateway DENIES all users by default for security.
Enter user IDs to create an allowlist, or leave empty
and you'll be asked about open access next.
[ Allowed user IDs (comma-separated): 
✓ Saved - only these users can interact with the bot.

For DMs, this is your user ID. You can set it later by typing /set-home in chat.
Home channel ID (for cron/notification delivery, or empty to set later with /set-home):
```

I'm choosing Telegram myself, so I created a new bot at @BotFather and provided its token to my Hermes agent.

Step 6 - start Hermes

That's it. You now have a running agent with an interactive CLI. Start giving it tasks or message it from Telegram.

```
HERMES-AGENT

Hermes Agent v0.9.0 (2026.4.13) · upstream 333cb825

Available Tools
browser: browser_back, browser_click, ...
clarify: clarify
code_execution: execute_code
cronjob: cronjob
delegation: delegate_task
file: patch, read_file, search_files, write_file
homeassistant: ha_call_service, ha_get_state, ...
image_gen: image_generate
(and 11 more toolsets...)


Available Skills
autonomous-ai-agents: claude-code, codex, hermes-agent, opencode
creative: architecture-diagram, ascii-art, ascii-video, e...
data-science: jupyter-live-kernel
devops: webhook-subscriptions
email: himalaya
gaming: minecraft-modpack-server, pokemon-player
general: dogfood
github: codebase-inspection, github-auth, github-code-r...
leisure: find-nearby
mcp: mcpserver, native-mcp
media: gif-search, heartmula, songsee, youtube-content
mlops: audiocraft-audio-generation, axolotl, clip, dsp...
note-taking: obsidian
productivity: google-workspace, linear, nano-pdf, notion, ocr...
red-teaming: godmode
research: arxiv, blogwatcher, llm-wiki, polymarket, resea...
smart-home: openhue
social-media: xitter
software-development: plan, requesting-code-review, subagent-driven-d...

29 tools · 76 skills · /help for commands
▲ 14 commits behind — run hermes update to update
```

We have installed the Hermes agent on our VPS server so it can trade and learn 24/7. Now

let's apply the full power of this framework to trading on Polymarket weather markets.

Hermes - weather trading agent (guide)



WEATHERBOT v3 - LIVE

USDC free: \$0.99 (rest is deployed in positions)
Positions: 10 open

City	Date	Bucket	Entry	Cost	EV
Chicago	04/17	82-83°F	\$0.225	\$2.00	+3.44
Dallas	04/17	84-85°F	\$0.340	\$1.82	+1.94
Seattle	04/17	52-53°F	\$0.110	\$1.37	+8.09
London	04/18	18°C	\$0.280	\$1.26	+2.57
Paris	04/17	20°C	\$0.195	\$1.02	+4.13
Seoul	04/17	17°C	\$0.190	\$0.95	+4.26
Ankara	04/17	19°C	\$0.065	\$0.77	+14.38
Tokyo	04/17	15°C	\$0.006	\$0.71	+165.67
Singapore	04/17	29°C	\$0.020	\$0.58	+49.00
Singapore	04/18	29°C	\$0.018	\$0.53	+56.14

Total deployed: ~\$11.01

8 of 10 positions resolve tomorrow (April 17 markets). London and Singapore D+2 resolve April 18. Bot is sleeping - next scan in ~50 min.

Now the payoff. We're going to deploy a production-grade, open-source Polymarket weather Hermes Agent so it runs 24/7, alerts you via Telegram, and self-calibrates over time.

This is the exact setup I use to run scans across 20 cities across 4 continents with three forecast sources feeding into Expected Value and Kelly Criterion position sizing.

we will use the logic of a weather trading bot made by @AlterEgo_eth

This step-by-step guide - just copy-paste these prompts to your Hermes agent. You don't need to code anything. Total time: ~30 minutes.

PROMPT 1 - Clone and setup

We will use an open-source bot made by AlterEgo to set up basic logic for our Hermes agent, so we send this command to our agent's CLI or TG bot.

PROMPT 2 - Create wallet

I recommend setting up a separate wallet for your weather bot. It's better to ask Hermes to set up a wallet for itself, then fund it.

Save the private key & wallet address that the bot will send you

PROMPT 3 - Fund the wallet (you do this yourself)

Send to your wallet address on Polygon network: \$ USDC.e - your trading capital (\$10 min, \$50 recommended), \$ POL - about 2 POL for gas (~\$1)

Once funded, tell your agent:

PROMPT 4 - Approve Polymarket contracts

Approve all contracts so the bot can use your funds for real trading

PROMPT 5 - Configure and connect weather API

Create an account at visualcrossing.com and get free API here, then send this prompt to your agent.

PROMPT 6 - Test scan

Run a test scan or a paper-trading mode before starting with real money.

PROMPT 7 - Start real trading

To start real trading, just send this command to your Hermes agent and witness the magic of a self-learning weather trading agent happening right in front of you.

As a result, you will see such a dashboard with your trade executions, and you may also ask your weather trading agent to send reports to your TG bot.

WEATHERBOT v3 – LIVE					
USDC free:		\$0.99 (rest is deployed in positions)			
Positions:		10 open			
City	Date	Bucket	Entry	Cost	EV
Chicago	04/17	82–83°F	\$0.225	\$2.00	+3.44
Dallas	04/17	84–85°F	\$0.340	\$1.82	+1.94
Seattle	04/17	52–53°F	\$0.110	\$1.37	+8.09
London	04/18	18°C	\$0.280	\$1.26	+2.57
Paris	04/17	20°C	\$0.195	\$1.02	+4.13
Seoul	04/17	17°C	\$0.190	\$0.95	+4.26
Ankara	04/17	19°C	\$0.065	\$0.77	+14.38
Tokyo	04/17	15°C	\$0.006	\$0.71	+165.67
Singapore	04/17	29°C	\$0.020	\$0.58	+49.00
Singapore	04/18	29°C	\$0.018	\$0.53	+56.14
Total deployed: ~\$11.01					

The result is a fully functional, self-learning weather trading agent that improves with each trade rather than following a fixed script.

Your next task is to give the bot enough trades to learn from and clear self-adjustment instructions based on that data.